



Application of DCGANs in drawing faces

Mohamed Dwidar

20100361

TA/ Mahmoud ElMorshedy

Network architecture:

I made the network having two models as instructed, one generator the other discriminator.

I also used mixed precision meaning the network would use 16-byte floats instead of 32-byte floats when it could because without it TensorFlow would throw an Out of VRam (OOM) error.

The GPU used for training: RTX 3060 laptop (6GB vram)

The generator:

Takes input shape of (100,)

Makes its way to the shape of (192,128,1) for a grayscale image

The data set was shaped to be compatible with the output of the network not the opposite

dense (Dense)	(None, 51200)	5120000
batch_normalization (Batch Normalization)	(None, 51200)	204800
p_re_lu (PReLU)	(None, 51200)	51200
reshape (Reshape)	(None, 16, 16, 200)	0
conv2d_transpose (Conv2DTranspose)	(None, 32, 32, 64)	320000
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 64)	256
p_re_lu_1 (PReLU)	(None, 32, 32, 64)	65536
conv2d_transpose_1 (Conv2DTranspose)	(None, 96, 64, 64)	102400
batch_normalization_2 (Batch Normalization)	(None, 96, 64, 64)	256
p_re_lu_2 (PReLU)	(None, 96, 64, 64)	393216
conv2d_transpose_2 (Conv2DTranspose)	(None, 192, 128, 1)	576

The discriminator:

This one uses a very normal CNN to detect real from fake nothing new here

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 192, 128, 64)	1664
p_re_lu_3 (PReLU)	(None, 192, 128, 64)	1572864
dropout (Dropout)	(None, 192, 128, 64)	0
conv2d_1 (Conv2D)	(None, 96, 64, 64)	102464
p_re_lu_4 (PReLU)	(None, 96, 64, 64)	393216
dropout_1 (Dropout)	(None, 96, 64, 64)	0
conv2d_2 (Conv2D)	(None, 32, 32, 32)	51232
p_re_lu_5 (PReLU)	(None, 32, 32, 32)	32768
dropout_2 (Dropout)	(None, 32, 32, 32)	0
flatten (Flatten)	(None, 32768)	0
dense_1 (Dense)	(None, 64)	2097216
p_re_lu_6 (PReLU)	(None, 64)	64
dense_2 (Dense)	(None, 32)	2080
dense_3 (Dense)	(None, 1)	33


The dataset:

The dataset was resized to 192*128 and gray scaled

The dataset size is 11GB so I loaded it as batches instead of loading the whole dataset to memory as it would cause the kernel to crash

The training results:

The loss chart:



The images:

